

Learning to Apply, Not to Memorize: Policy-in-Context Fine-Tuning Generalizes to Unseen Organizational Policies

Anonymous ACL submission

Abstract

Organizations want AI systems that make decisions *their* way: following the written policy, citing the rule that applied, and deferring exactly where the policy grants discretion. The dominant fine-tuning recipe internalizes a specific policy into model weights, which breaks the moment the policy changes. We study the opposite framing: supply the policy *in context* at inference time, and fine-tune a small model on the **skill of applying whatever policy it is given**. On a corpus of 34 organizational decision cases (3,400 labeled decisions generated by a deterministic rules engine over declarative ontologies), a QLoRA-tuned 7B model reaches **72.9% exact match** (verdict, routing, parameters, and a *correct* cited rule) on policies never seen in training—evaluated leave-case-out so every policy is unseen exactly once—versus 15.0% for the untrained model under identical serving (cluster-bootstrap 95% CI [64.7, 80.9] vs. [3.3, 26.7]). Three controls pin down what was learned: on **counterfactual policies** (approve/reject flipped, evaluated only where the flip changes the answer, so prior knowledge scores zero) the tuned model scores **88.0%**; a 3-shot in-context-learning baseline reaches only 50.0%; and training-seed variance is ± 4.8 pp. A negative control bounds the recipe: the *same* recipe on an exact arithmetic derivation task scores 0/6—even with chain-of-thought targets, and even for a frontier model prompted identically without tools—confirming that derivations belong outside inference: not merely behind a per-run tool call, but compiled to code so repeat runs carry no LLM at all. Because the pipeline compiles from a work session the organization already paid for, the frontier-token break-even is the *second run*.

1 Introduction

When an employee asks “can I give this customer a 12% discount?”, the organization does not want

an oracle’s opinion. It wants *its own policy applied*: the ordered rules checked top-down, the first matching rule cited, the case routed to the approver the policy names, and the band the policy deliberately left open (“5–10%: AI may recommend, team lead approves”)—and only that band—filled by a recommendation. Three properties shape the ML problem:

1. **Policies change faster than models.** A model that memorized policy P is wrong the day P becomes P' . The policy must live *outside* the weights.
2. **Judgments must be auditable.** “Approve” is not an answer; “approve, route to director, citing strategic_retention_band whose condition holds on this record” is.
3. **Discretion is declared, not inferred.** The policy says where the model may recommend; confidence thresholds do not.

Prior work internalizes policies into weights (Wang et al., 2025), benchmarks agents given domain policy guidelines (Yao et al., 2024; Zhou et al., 2024), or measures robustness to rule updates (Diao et al., 2025). What is missing is the combination deployment actually needs: **fine-tune a small, locally servable model on the transferable skill of applying an arbitrary in-context policy, and measure transfer to policies never seen in training.**

Organizational policy decisions are our testbed, but the object of study is more general: a **content-in-context, skill-in-weights** decomposition. The same shape recurs wherever an organization’s *documents* define correct behavior and change on the organization’s schedule—support runbooks, escalation matrices, grading rubrics, editorial style guides, compliance checklists. We use declarative decision policies because they are the cleanest instance of the class that still admits a deterministic

grader: the skill (ordered evaluation, first-match, grounded citation) is non-trivial, while correctness is checkable to the byte. Whether the recipe transfers to messier members of the class is an open question we scope explicitly (§6).

Our contributions:

- **Task & corpus (§3):** 34 decision cases across 10 organizational functions, each declared as an ontology plus ordered rules with explicit discretion bands, compiled by a deterministic engine into 3,400 labeled decisions with cited rules and grounded conditions.
- **A strict deterministic grader (§3.3):** exact match on verdict, route, parameters *and* cited rule, plus a check that the cited rule’s condition actually holds on the record.
- **The transfer result (§5.1):** leave-case-out over all 34 policies, 72.9% exact [64.7, 80.9] vs. 15.0% raw under identical bf16 greedy serving; ICL ladder 15.0 → 50.0 → 72.9.
- **The counterfactual control (§5.2):** 88.0% on flipped policies where priors score zero by construction.
- **Negative control and placement economics (§5.3–5.4):** the same recipe on byte-exact arithmetic fails (0/6), survives a CoT-target ablation, and is matched by a tool-less frontier model failing the same gate—justifying compilation to code over per-run tool calls, with a measured second-run break-even.

2 Related Work

Policy internalization vs. policy-in-context. Multimodal Policy Internalization (Wang et al., 2025) trains policies *into* model parameters so the policy need not be in context at inference—the opposite maintenance bet from ours. Its “Policy Override” setting supplies an updated policy in context at inference as a *generalization probe* of the internalized model; we make in-context application the training objective itself. Policy Reasoning Traces (Imperial and Madabushi, 2025) improves compliance assessment on specific policies such as HIPAA and GDPR with per-policy instance-level train/test splits; cross-policy transfer appears only as an auxiliary analysis, not as the held-out-policy protocol we use. GuideBench (Diao et al., 2025) benchmarks robustness to rule updates.

Policy-following benchmarks. τ -bench (Yao et al., 2024) evaluates state-of-the-art agents given domain policy guidelines; RuleArena’s (Zhou et al., 2024) findings distinguish failures of rule identification from failures of the computations the rules require—directly anticipating our §5.3; CRMArena-Pro (Huang et al., 2025) assesses agents across enterprise scenarios. A contemporaneous enterprise benchmark (Yu et al., 2026) independently builds rule-engine-judged decisions; it is evaluation-only and has no discretion bands. DMN (Object Management Group, 2024) has long standardized ordered decision rules with FIRST-hit semantics; our DSL is deliberately in that family—the ML contribution is the transfer measurement, not the rule language.

SFT and arithmetic. Faith and Fate (Dziri et al., 2023) shows Transformers fail at compositional multi-step arithmetic; Chu et al. (2025) shows SFT tends to memorize and struggles out-of-distribution; PAL (Gao et al., 2022) offloads the solution step to a Python interpreter. Our contribution is the *pairing*: the same recipe, corpus scale and gate that fail on arithmetic succeed on rule application—and a tool-less frontier model fails the same arithmetic gate, sharpening “small models can’t compute” into “inference can’t compute.”

Small-model deployment. Belcak et al. (2025) argue small models are the future of agentic AI and outline an LLM-to-SLM conversion algorithm; we supply a trained-skill recipe and transfer measurement for one high-value slot (policy application) and a deterministic gate that even frontier models fail for another (derivation), yielding a placement rule rather than a preference.

3 Task and Corpus

3.1 Decision cases

A *case* declares one recurring organizational decision: the deciding role, an **ontology** (entities, attributes, relations), a feature space, **ordered rules** (when condition → verdict, route, parameters, rationale), optional **discretion bands** (defer: slm_recommend with the band’s limits as parameters), and a fallback. The corpus holds **34 cases across 10 functions** (sales, finance, support, HR, procurement, legal, IT operations, marketing, logistics, security). A deterministic engine samples records from each case’s feature space and applies first-match semantics to produce **100 labeled deci-**

sions per case (3,400 total), each with the applied rule’s name, its condition, and a rationale. Generation is seeded and byte-reproducible.

3.2 Policy-in-context format

Per instance the model receives: a system turn naming the organization, role and decision; the case’s ontology and its *full ordered rule list* (YAML, verbatim); and the record. It must output one JSON object {*verdict*, *route*, *params*, *cited_rule*, *rationale*}. Nothing about any specific policy is in the weights; a rule edit changes behavior with no retraining.

3.3 Grading

An answer is **exact** iff verdict, route and parameters equal the engine’s decision *and* the cited rule is the first-matching rule *and* that rule’s condition actually holds on the record (the last check fails fabricated grounds even when the verdict is guessed right). A **relaxed** variant accepts a cited rule that is not first-match but whose condition genuinely holds with all other fields correct. We also report verdict-only accuracy. Because held-out items cluster by policy, all confidence intervals are **case-cluster bootstrap** (10,000 resamples over cases).

3.4 Honest scope

The corpus is synthetic and self-graded: the engine that labels is the engine that grades, the policies are already machine-readable, and rule-condition field names match record keys—the entity-alignment difficulty of prose policies is absent, as are noisy fields and annotator disagreement. We claim transfer of *policy application on declarative policies*, a deliberately clean substrate; §6 discusses the path to prose policies.

4 Experimental Setup

Model & training. Qwen2.5-7B-Instruct, QLoRA ($r=16$, $\alpha=32$, dropout 0.05, all linear projections), completion-only loss, 2 epochs, LR 10^{-4} cosine, bf16 compute on one RTX 4090 (≈ 12 min/run). Training data: 28 cases $\times$ 40 instances (1,120 rows); 6 cases held out *entirely*. The environment is pinned and every run writes a manifest (seed, epochs, data and adapter SHA-256, package versions).

Serving & fairness. All models—tuned *and* raw baselines—are evaluated through the same bf16

model (identical serving)	exact	95% CI
raw 3B, zero-shot	11.7%	[0.0, 25.0]
raw 7B, zero-shot	15.0%	[3.3, 26.7]
raw 7B, 3-shot ICL	50.0%	[28.3, 71.7]
tuned 7B (pinned split)	68.3%	[53.3, 83.3]
tuned 7B (LOCO)	72.9%	[64.7, 80.9]

Table 1: Unseen-policy exact match. LOCO holds every one of the 34 policies out exactly once.

server with greedy decoding. (Re-running raw baselines under identical serving moved raw 7B from 16.7% to 15.0%; the quantization confound does not explain the gap.)

Splits. (a) *Pinned*: 6 held-out cases $\times$ 10 instances, plus 56 seen-policy/unseen-instance items. (b) *Leave-case-out (LOCO)*: 6 folds re-trained with the same recipe so every one of the 34 policies is unseen exactly once (340 items, 34 clusters).

5 Results

5.1 Transfer to unseen policies

LOCO fold accuracies range 65.0–80.0%: the pinned-split result is not an artifact of one partition (Table 1, Figure 1a). Across 3 training seeds, unseen exact is $62.8\% \pm 4.8\text{pp}$ —run variance is real and reported; no seed approaches the raw baseline. In-context examples help substantially (15 $\rightarrow$ 50%) but fine-tuning adds $\sim 20\text{pp}$ over ICL while removing per-request example tokens. Per-policy scores range 1/10–10/10 under LOCO (Figure 2); policies whose conditions embed date/limit arithmetic (leave, overtime, payment extension) sit at the bottom, foreshadowing §5.3.

5.2 What was learned: the counterfactual control

We flip approve $\leftrightarrow$ reject in every rule (and the fall-back) of the held-out cases and evaluate **only on records where the flipped policy decides differently from the original**—on this subset, answering from priors or from the original policy scores zero by construction. The tuned model scores **88.0%** (44/50); raw 7B scores 48.0% (Figure 1b). Combined with §5.1, this is direct evidence that fine-tuning taught *application of the in-context policy*, not policy content and not world priors.




Figure 1: (a) Unseen-policy exact match with case-cluster bootstrap 95% CIs. (b) The counterfactual control: on records where the flipped policy decides differently, priors score zero by construction.




Figure 2: Per-policy transfer under LOCO; every policy is unseen exactly once ($n=10$ each).

arm	gate	field acc.
SFT 7B (plain targets)	0/6	—
SFT 7B (CoT targets)	0/6	86.4%
frontier, same prompt, no tools	0/6	93.6%

Table 2: The arithmetic negative control. Both tuned and frontier arms fail almost solely on two long-division fields.

5.3 Negative control: the same recipe cannot learn computation

We ran the identical recipe (same trainer, gate, corpus scale: 300 examples) on a *derivation* task from the same product family: regenerate a customer’s renewal-pricing files (22 numeric fields) byte-exactly for held-out customers, judged by a deterministic gate.

This is a *control*, not a *discovery*—routing arithmetic to tools is settled practice. Its role: (i) it

closes the “train harder” objection (the CoT-target ablation still fails; the frontier fails the same gate on the same two long-division fields while getting every band and discount rule right; the recorded agent originally passed because it computed with tools); (ii) field-level scores show what the exact gate hides—inference gets 20–21 of 22 fields, failing precisely where computation, not rule-reading, is required; (iii) it motivates going beyond the standard remedy: a per-run **tool call keeps the LLM in the loop on every execution** (re-deciding which tool to invoke, nondeterministically, with orchestration tokens), whereas compiling the step to the **code tier makes that decision once**, at compile time. The placement rule: derivations $\rightarrow$ code; policy judgments $\rightarrow$ trained SLM behind the deterministic gate; declared discretion bands $\rightarrow$ recommendation + human approver.

5.4 Placement economics: break-even at the second run

Measured on the renewal pipeline in unique tokens (not double-counting the cumulative context an agent resends each turn): the recorded agent session costs 23,614 frontier tokens per run; after compilation and SLM promotion a run consumes **0 frontier tokens** (4,208 tokens on a local model at \$0 marginal cost, $7.4\times$ faster, outputs reproduced 7/7 under a deterministic benchmark). Because the compiler’s input is a work session the organization had to run once anyway, the frontier-token **break-even is the second run**—compared with the ≈ 17 -transaction break-even reported for optimization-heavy compilation (Trooskens et al.,

2026), where the compiler spends LLM tokens searching. Escalations for genuinely new parameter values each cost one frontier call, are cached with upstream-output fingerprints, and amortize to zero on repeats. The hidden denominator is the human cost of declaring the ontology and policy; the pipeline pays off in proportion to a decision’s repetition frequency—exactly the regime organizations automate first.

6 Discussion

Why not internalize? A weights-internalized policy answers fast but ages instantly and cannot explain a rule edit. Policy-in-context inverts the maintenance economics: the policy file is the deployment artifact; retraining is only needed when the *skill*, not the policy, drifts.

Beyond policy decisions. Nothing in the recipe is specific to approval workflows: training rows are (document, record, grounded structured answer) triples, and the counterfactual control generalizes to any document class whose semantics can be programmatically perturbed. Runbooks and rubrics are the nearest neighbors; each needs its own deterministic grader, which is the actual bottleneck for widening the claim. This paper is also one tier of a larger compiled-pipeline system—recorded agent sessions lowered to code/rule/cache tiers, with the trained-SLM tier occupying the judgment slot behind a promotion gate—whose end-to-end economics (§5.4) motivate but do not depend on the transfer result.

Path to prose policies. Our substrate removes entity alignment (rule field names match record keys). The natural next experiment keeps the trained applier fixed and inserts a policy-compilation step (prose → declarative rules) in front—measuring how much of the 72.9% survives noisy compilation, and whether the counterfactual control still holds.

Limitations

Labels and grades come from the same engine; we mitigated with the grounds-hold check, the relaxed metric, the counterfactual control and full pinning, but human-annotated decisions with inter-annotator agreement remain future work. Scale: 34 policies, one 7B base model, one language pair (Korean prompts, mixed-language fields); cluster CIs are honest about this—[64.7, 80.9] is the

claim, not 72.9%. Policies are already machine-readable; prose-policy entity alignment is untested. The frontier control uses one frontier model, one prompt format, without tools by construction.

Reproducibility

Corpus generator, trainer (pinned; per-run manifests), evaluation code, all raw results and this draft live in one repository; every table is regenerable from a seeded generator and stored model outputs.

References

- Peter Belcak et al. 2025. Small language models are the future of agentic ai. ArXiv:2506.02153.
- Tianzhe Chu et al. 2025. Sft memorizes, rl generalizes: A comparative study of foundation model post-training. ArXiv:2501.17161.
- Lingxiao Diao, Xinyue Xu, Wanxuan Sun, Cheng Yang, and Zhuosheng Zhang. 2025. Guidebench: Benchmarking domain-oriented guideline following for llm agents. ArXiv:2505.11368.
- Nouha Dziri et al. 2023. Faith and fate: Limits of transformers on compositionality. ArXiv:2305.18654, NeurIPS 2023.
- Luyu Gao et al. 2022. Pal: Program-aided language models. ArXiv:2211.10435.
- Kung-Hsiang Huang, Akshara Prabhakar, Onkar Thorat, Divyansh Agarwal, Prafulla Kumar Choubey, Yixin Mao, Silvio Savarese, Caiming Xiong, and Chien-Sheng Wu. 2025. Crmarena-pro: Holistic assessment of llm agents across diverse business scenarios and interactions. ArXiv:2505.18878.
- Joseph Marvin Imperial and Harish Tayyar Madabushi. 2025. Scaling policy compliance assessment in language models with policy reasoning traces. ArXiv:2509.23291.
- Object Management Group. 2024. Decision model and notation (dmn) v1.5. OMG standard.
- Geert Trooskens, Aaron Karlsberg, Anmol Sharma, Lamara De Brouwer, Max Van Puyvelde, Matthew Young, John Thickstun, Gil Alterovitz, and Walter A. De Brouwer. 2026. Compiled ai: Deterministic code generation for llm-based workflow automation. ArXiv:2604.05150.
- Zhenhailong Wang, Jiateng Liu, Amin Fazel, Ritesh Sarkhel, Xing Fan, Xiang Li, Chenlei Guo, Heng Ji, and Ruhi Sarikaya. 2025. Multimodal policy internalization for conversational agents. ArXiv:2510.09474.

- Shunyu Yao, Noah Shinn, Pedram Razavi, and Karthik Narasimhan. 2024. τ -bench: A benchmark for tool-agent-user interaction in real-world domains. ArXiv:2406.12045.
- Tao Yu, Hao Wang, Changyu Li, et al. 2026. Beyond the all-in-one agent: Benchmarking role-specialized multi-agent collaboration in enterprise workflows. ArXiv:2605.08761.
- Ruiwen Zhou, Wenyue Hua, Liangming Pan, Sitao Cheng, Xiaobao Wu, En Yu, and William Yang Wang. 2024. Rulearena: A benchmark for rule-guided reasoning with llms in real-world scenarios. ArXiv:2412.08972.